

MERIVALE HIGH SCHOOL  
INTERNATIONAL BACCALAUREATE PROGRAMME

---

# **Investigating the Impact of Various Feature Extraction Algorithms on Performance in Automatic Speech Recognition Systems**

*To What Extent Do Mel-Frequency Cepstrum Coefficients Outperform Other  
Speech Feature Extraction Algorithms in Speech Recognition Neural Networks?*

---

*Student:*  
Daniel LU

*Supervisor:*  
Courtney EDWARDS

November 10, 2022

3662 Words

# Contents

# List of Abbreviations

<b>ASR</b>	<b>Automatic Speech Recognition</b>
<b>LPC</b>	<b>Linear Predictive Coding</b>
<b>DFT</b>	<b>Discrete Fourier Transform</b>
<b>FFT</b>	<b>Fast Fourier Transform</b>
<b>IFT</b>	<b>Inverse Fourier Transform</b>
<b>HMM</b>	<b>Hidden Markov Model</b>
<b>CTC</b>	<b>Connectionist Temporal Classification</b>
<b>DDC</b>	<b>Directed Dialogue Conversations</b>
<b>NLC</b>	<b>Natural Language Conversations</b>
<b>NLP</b>	<b>Natural Language Processing</b>
<b>DWT</b>	<b>Discrete Wavelet Transform</b>
<b>PLP</b>	<b>Perceptual Linear Prediction</b>



## Chapter 1

# Introduction

### 1.1 Abstract

Automatic speech recognition (ASR) is the ability of a machine to recognize language from human speech and convert it into written text. Such a task is difficult due to the sheer complexity of any spoken language, as well as the varying and unpredictable nature of different people and their speaking habits. Even big-name companies such as Alphabet, Microsoft, and Baidu, who pride themselves in their sophisticated multi-million dollar speech recognition systems, have been countlessly ridiculed for their misinterpretations. However, due to the endless number of possible and practical applications of ASR technology in our developing world, major corporations still aspire to develop an accurate yet optimized speech recognition system for their products and services. For instance, Baidu spent a total of \$4.7 billion dollars on research and development, of which a large portion headed into developing their voice assistant, DuerOS. Hence, the implementation of ASR can improve the ergonomics of any task that requires human interaction, such as personal assistants, smart healthcare and household appliances, and media transcription to support people with disabilities or language barriers.

Significant development for ASR began in the 50s, with the innovation of Bell's Audrey and IBM's Shoebox. These systems fell into a category of ASR systems that recognized individual words to an isolated vocabulary using manually set parameters for phoneme recognition. By the end of the decade, ASR technology could distinguish words with a maximum of four vowels and nine consonants and by the 80s, systems could recognize up to 3000 words, which is around the vocabulary of a six-year-old child. It wasn't until the late 90s that researchers transitioned outwards of phoneme detection to Hidden Markov Models (HMM) and focused on natural language processing. Deep learning concepts and technologies emerged and were investigated as a solid alternative to HMMs, to the extent to which it is prominently used today.

However, a plausible deep learning system that can recognize continuous speech within a low margin of error is very strenuous to build. A well-designed system can consist of multiple complex layers that can be very complicated to integrate into one another. Nowadays, large tech companies boast an accuracy of almost 95

## 1.2 Rationale

This report aims to investigate the extent to which Mel-Frequency Cepstrum Coefficients outperform other feature extraction methods in automatic speech recognition systems. I will assess the performances of MFCCs, Discrete Wavelet Transforms (DWT), Linear Predictive Coding (LPC), and Perceptual Linear Prediction (PLP), as well as the absence of a feature extraction layer, processed through identical recurrent neural network architectures. From this investigation, I will better understand the effectiveness and limitations of different speech feature extraction methods and inform my forthcoming decisions in ASR neural networks.

## Chapter 2

# Background

### 2.1 Automatic Speech Recognition


Automatic speech recognition models can be split into two broad categories of complexity. First of all, there is single-word classification to an isolated vocabulary where directed dialogue conversations are analyzed; its functionality can be seen in early speech recognition efforts such as IBM's Shoebox. The other variant is continuous speech recognition, where natural language conversations (NLCs) are processed and transcribed to a vocabulary of around 20,000 to 100,000 words. Many technologies today utilize this system, including video subtitles, voice searching, and customer services at call centres. Directed dialogue classification models can be fairly simple to understand and produce through various feed-forward networks, but natural language networks require the use of recurrent neural networks.

### 2.2 Recurrent Neural Networks for ASR

Recurrent neural networks are a variant of deep learning neural networks that use time-sequential data without a fixed shape. These systems are inspired by the biological composition of neurons in the brain, as humans have learned long ago to adapt their innovations from nature. Each neuron, out of around 100 million, is interconnected to many other neurons and passes a signal via its axon to receiving axon terminals belonging to other neurons. Contrastingly, in an artificial neural network, a combination of weights and biases stacked in layers are constantly optimized through gradient descent to produce a matrix of probabilities that highlight the best possible output the network can generate at its current state (for the case of classification architectures). In each cell, a dot product of the input matrix and the cell's weight matrix is done and summed with the cell's bias matrix to create an output matrix of values, fed into all of the nodes in the next layer. At first, these matrices use senseless, randomly generated weights and biases that result in a completely random output. However, as the model calculates its loss from an evaluation of the net deviation of the predicted output from the ground truth, a gradient is formulated that adjusts each layer's parameters closer to what would create an output with a lower loss. The higher loss an output receives, the greater magnitude of the gradient is passed to backpropagation and the more each node is changed, with the

vice versa applying as well.

In a typical multi-layer deep learning neural network, the input and output shape of the model are fixed and have to be set. Since the length of the NLC audio input is indeterminate in this investigation, recurrent neural networks can be used to process continuous speech data that last an indeterminate length of time. RNNs accomplish this task by accepting an input for each time step into a recurrent cell and feeding the cell's generated output to the succeeding cell through a hidden layer, doing so until the input data is completely used. This results in the coupled nature of a recurrent neural network that enables it to find relationships in continuous data, namely in stock market prediction, and influence future outputs. The framework of a standard RNN cell is laid out below (Figure 2.1).



$$h_t = \tanh(W \left( \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix} \right))$$

FIGURE 2.1: The internal composition of a hyperbolic tan RNN cell

The distribution of input data and output data can vary from network to network and is ultimately based on the shape of the input data and the desired shape of the output data. For example, a music generation RNN could take in a dataset of Vivaldi's works as an input and output an infinite length of completely original Vivaldi-styled music: this model would be classified as a one-to-many RNN structure. The following table (Table 2.1) describes four main types of RNN structures that satisfy most machine learning purposes.

Due to the continuous nature of NLC, multiple inputs as the audio data and its respective outputs as phonemes, letters, or words are present. Thus, a many-to-many structure will be used as the framework for this investigation's experimentation.

However, conventional recurrent neural networks have two major problems that need to be considered: the vanishing and exploding gradient problem, and the long-term dependency on sound signals. The vanishing and exploding gradient problem appear during the backpropagation stage of the network, where a gradient propagates through each




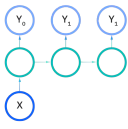
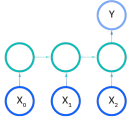
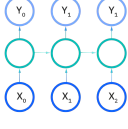
Architecture	Characteristics	Application	Diagram
One-to-One	Single input Single output	Image classification	
One-to-Many	Single input Sequence of outputs	Image captioning	
Many-to-One	Sequence of inputs Single output	Sentiment analysis	
Many-to-Many	Sequence of inputs Sequence of outputs	Language translation	

TABLE 2.1: Four types of recurrent neural network structures


cell and updates the weights and biases of each node in order for the network to learn. The problem arises when the number of layers starts to increase and the gradients start to travel and multiplied through more and more layers; when a gradient with a slight deviation from  $y=x$  is returned to the first node during backpropagation, it begins to increase or decrease exponentially (dependent on its sign) as it progresses through each cell. Correspondingly, as it reaches near the end of backpropagation, the gradient is either too minuscule to the point where the end weight is scarcely updated at all, or too large where the weight cannot be optimized properly.

Going on, sound signals can take up a significant portion of input data, dependent on sample frequency, as a one-second audio clip of someone dictating a word can take up to 16,000 time steps of data. The extreme short-term memory of a simple hyperbolic tangent recurrent network cannot remember what sound signal was inputted into the cell some 5,000 iterations earlier, to the extent it is significant towards producing the output. This is detrimental towards the ultimate performance of the neural network, as a large segment of a model's decision on a phoneme, character, or word depends on contextual data beforehand. For example, a model could have a hard time discriminating between the letter "k" and "q" with no context. However, the prerequisite knowledge that it was preceded by a "loo" will likely increase the probability that the following letter is "k". These two setbacks of recurrent neural networks introduced a revolutionary recurrent cell concept introduced in 1997 and prominently used today, termed the Long Short-Term Memory cell.

## 2.3 Long Short-Term Memory

Long Short-Term Memory cells counter the two problems above by incorporating an extra hidden memory state that propagates throughout recurrences of the network. The addition of a memory state that can remember important specifics from many recurrences beforehand and influence the current output. Contrary to the perception-based framework of traditional neural networks, an LSTM cell is much more convoluted and contains substantially more weights and biases in the form of specialized gates (Figure 2.2). The combination of these properties provides the fundamental framework that conceives the LSTM network's ability to recognize context-sensitive information and is thus the reason why these networks are the standard for building reliable speech recognition systems.

A simple speech recognition system is usually composed of two different components; a feature extraction layer and the neural network itself. The purpose of the feature extraction layer is to process the raw input data into a more usable and space-efficient form. In the case of ASR, input data is usually in the form of waveform data embedded in .mp3 or .wav files. This one-dimensional waveform data has several limitations due to having less useful features than a two-dimensional format of audio signals, such as spectrograms. Spectrograms are a visual representation of the amplitudinal spectrum of



$$\begin{pmatrix} i_{gate} \\ f_{gate} \\ o_{gate} \\ g_{gate} \end{pmatrix} = \begin{pmatrix} \sigma \\ \sigma \\ \sigma \\ \tanh \end{pmatrix} W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix}$$

$$\begin{aligned} c_t &= f \cdot c_{t-1} + i \cdot g \\ h_t &= o \cdot \tanh(c_t) \end{aligned}$$

FIGURE 2.2: The internal composition of an LSTM cell

frequencies of an audio signal with respect to time. These isolated frequency amplitudes are also called the Discrete Fourier Transform (DFT) of an input and they can be extracted via the Fast Fourier Transform (FFT) algorithm, laid out below.

DFTs are frequently used in convolutional neural networks due to their property to be represented as an image and are the reliable standard in terms of extracting audio features from waveform data. In a DFT, the x-axis refers to time and the y-axis refers to the frequency that's played at a certain point in time, with a lighter colour indicating that the amplitude is larger with regard to magnitude (Figure 2.3).




FIGURE 2.3: A spectrogram of a person saying the phrase, "Taco Bell"

## 2.4 Feature Extraction Algorithms

However, spectrograms are not very well suited nor optimized for visualizing human speech. For instance, a human voice at a conversational level will normally not exceed a frequency of 500 Hz and fall below 60 Hz and it's not necessary to include non-phonetically vital properties of speech signals in the input data. Therefore, Mel Frequency Cepstral Coefficients (MFCC) intended to artificially replicate the human hearing system with the assumption that the human ear is a reliable speech recognition system in itself. A precise logarithm function is used to retain the phonetically vital properties of speech signals and plot the result on the Mel scale (below).

$$mel(f) = 2595 \times \log_{10}(1 + f/700)$$

MFCCs are created first by passing spectrogram data through a precise logarithm function, dubbed the Mel-filter bank, to retain the phonetically vital properties of speech signals while also reducing the number of unwanted frequencies in our input data. The result is plotted to a Mel spectrum and applied to the inverse discrete cosine transform, in which it is finally converted into a cepstrum, the inverse of a spectrum. The result is a set of 10-21 coefficients that ultimately present the main speech features of the original waveform data. MFCC data can also be taken as the derivative of and provide further feature data to feed into the neural network.

Other speech feature extraction functions that will be investigated are DWT, LPC, and PLP. They are each unique in their own manner but are also very closely related to each other with reference to the algorithms used to create them. Discrete Wavelet Transforms decompose a signal into a set of mutually orthogonal wavelet basis functions. It uses the Short-Time Fourier Transform, a variant of the FFT, which carries out a Fourier transform on a signal split into small windows of fixed duration instead of a specific time instance. Linear Predictive Coding generates coefficients that, similar to MFCCs, reflect the characteristics of a simplified vocal tract model, overall compressing the signal alongside the process. As the name suggests, the method is competent at predicting the future of a random process given past observations. Finally, Perceptual Linear Prediction claims to improve speaker-independent recognition performance and outperform LPCs in speech recognition. It first warps the spectrogram along the Bark frequency scale (Figure 2.4), convolves it with the power spectrum, and performs a cubic-root amplitude compression, ultimately creating the linear prediction model.




FIGURE 2.4: The Bark scale, visualized on a graph comparing Bark value and frequency

## Chapter 3

# Experiment Materials & Methodology

In this paper, heavy importance is placed on primary experimental observations and data, as well as a complete understanding of every operation that is carried out. For those reasons, deep learning APIs like Keras or Torch were not implemented, though this might have an effect on the overall performance of my neural network due to the great amount of optimization such open-source projects have. Nevertheless, this doesn't have an impact on the results of my investigation, as the network will remain a control variable throughout all trials. The dependent variables are the speech feature extraction algorithms (MFCC, DWT, LPC, and PLP) and the background noise of each sample (none, white noise, and miscellaneous/random). To properly appraise the extent to which MFCCs can extract speech features, background noise must be a variable due to being a known downside to them. The independent variables are the accuracy and time allotted to each network when processing the test dataset.

### 3.1 Datasets Used

I used "Speech Commands: A Dataset for Limited-Vocabulary Speech Recognition" published by Pete Warden to train my neural network. The dataset contains 6500 sets of one-word dialogue paired with its transcription as the true output. I later altered this dataset by concatenating copies of the audio data with various types of background noise using Librosa.

However, as I plan to use continuous speech recognition neural network architectures, using this dataset is only a simplistic and intuitive demonstration of what the model is capable of doing, in addition to an effort by myself to not be too ambitious and over-extend my experimentation for this report. Further developments of this model would feature and employ other open-source datasets that include multi-word speech data, such as podcasts or court hearings. I hope to alter this later on.

## 3.2 Pre-Processing

The dataset provided the sound files in .wav format, which made it very simple to extract its waveform data. I used ARSS's audio analysis program to perform an FFT on the dataset and obtain the respective DFT spectrograms. Going on, I used Librosa to retrieve the spectrograms' speech feature data to ultimately feed into the network.

## 3.3 Network Architecture

This report uses an end-to-end recurrent phoneme-based neural network consisting of a layer of LSTM cells and a final decoding node to process and display the output. The LSTM model consists of 256 hidden units that accept an input that's of shape (3, 26) and output a one-hot matrix of shape (46, 1) at each time step. The output is a probability matrix of all the 44 English phonemes, following ARPABET conventions, that the model predicts the audio signal refers to at that point in time, as well as a space and a "blank" character that are incorporated to separate concurrent letters in the combined output. I choose to use a phoneme classification system over character-based or word-based since I believe there is more flexibility and accuracy established when you dissect a sentence into the fundamental building blocks of spoken language and build from the ground up once the phonemes are selected. A working result of this model would output a sequence of slurred letters that, once decoded with Connectionist Temporal Classification (CTC) loss, would output a word; the model could output "ww\_eee\_ee\_p", which suggests that the original audio input is of a person saying "weep". After the fact, these phonemes are then fed into an English language model to convert the outputted phonemes into written text.

The model was trained on 5000 pieces of training data, 10 epochs of such, with a learning rate of 0.01 and evaluated by three validation datasets, for each of the three trials, that held 500 pieces of data to determine the accuracy of each model.

## 3.4 Experimental Process

Before I elaborate on my experimental process, it is essential to clarify some notable semantics of this investigation. The general performance basis of a model is assessed on a synthesis of its accuracy and its efficiency, though with a greater emphasis on accuracy. Whether an ASR system is successful is determined by whether its accuracy is greater than 50% when recognizing NLC and whether its processing latency is less than five seconds. The accuracy of a model is calculated appertaining to its loss from stochastic CTC and the efficiency of a model is gauged on a comparison of its estimated time complexity and a measured difference of time.

I will hold three trials for each set of independent variables to answer my research question. The performance of raw spectrograms, MFCCs, DWTs, LPCs, and PLPs will be assessed against each other in three different dataset variants: with no background noise, with white noise, and with other miscellaneous background noises, amounting to a total of 45 trials.



## Chapter 4

# Experimental Results

### 4.1 Tabular Presentation of Data

The tables below show the experimental results of all the trials, with each table presenting the performance of each feature extraction algorithm in varying types of noise.

Algorithm	Noise	Trial	Accuracy (%)	Processing Time (s)
FFT Spectrogram	None	<i>Trial 1</i>	0.937	4.36
		<i>Trial 2</i>	0.938	4.80
		<i>Trial 3</i>	0.921	4.72
		Average	0.932	4.63
	White	<i>Trial 1</i>	0.862	4.29
		<i>Trial 2</i>	0.887	4.63
		<i>Trial 3</i>	0.894	4.73
		Average	0.881	4.55
	Misc.	<i>Trial 1</i>	0.686	4.85
		<i>Trial 2</i>	0.621	4.90
		<i>Trial 3</i>	0.684	4.35
		Average	0.664	4.70
Mel Spectrogram	None	<i>Trial 1</i>	0.927	4.53
		<i>Trial 2</i>	0.948	4.88
		<i>Trial 3</i>	0.950	4.68
		Average	0.941	4.69
	White	<i>Trial 1</i>	0.767	4.41
		<i>Trial 2</i>	0.793	4.49
		<i>Trial 3</i>	0.788	4.94
		Average	0.782	4.61
	Misc.	<i>Trial 1</i>	0.626	934.93
		<i>Trial 2</i>	0.613	4.62
		<i>Trial 3</i>	0.640	4.46
		Average	0.626	4.67
	None	<i>Trial 1</i>	0.866	4.27
		<i>Trial 2</i>	0.855	4.55

Table 4.1 continued from previous page

Algorithm	Noise	Trials	Accuracy (%)	Processing Time (s)
	White	<i>Trial 3</i>	0.869	4.19
		Average	0.863	4.34
		<i>Trial 1</i>	0.783	4.13
		<i>Trial 2</i>	0.758	4.53
		<i>Trial 3</i>	0.750	4.43
		Average	0.764	4.36
	Misc.	<i>Trial 1</i>	0.579	4.28
		<i>Trial 2</i>	0.585	4.41
		<i>Trial 3</i>	0.610	4.41
		Average	0.592	4.37
DWT	None	<i>Trial 1</i>	0.897	4.67
		<i>Trial 2</i>	0.841	4.46
		<i>Trial 3</i>	0.874	4.44
		Average	0.870	4.52
	White	<i>Trial 1</i>	0.873	4.43
		<i>Trial 2</i>	0.834	4.42
		<i>Trial 3</i>	0.869	4.37
		Average	0.858	4.40
	Misc.	<i>Trial 1</i>	0.756	4.53
		<i>Trial 2</i>	0.734	4.48
		<i>Trial 3</i>	0.752	4.31
		Average	0.747	4.44
LPC	None	<i>Trial 1</i>	0.882	4.61
		<i>Trial 2</i>	0.833	4.91
		<i>Trial 3</i>	0.840	4.39
		Average	0.852	4.64
	White	<i>Trial 1</i>	0.849	4.53
		<i>Trial 2</i>	0.826	4.70
		<i>Trial 3</i>	0.857	4.57
		Average	0.844	4.6
	Misc.	<i>Trial 1</i>	0.756	4.61
		<i>Trial 2</i>	0.694	5.10
		<i>Trial 3</i>	0.732	4.79
		Average	0.747	4.83
PLP	White	<i>Trial 1</i>	0.832	4.55
		<i>Trial 2</i>	0.839	4.76
		<i>Trial 3</i>	0.861	4.77
		Average	0.843	4.73
		<i>Trial 1</i>	0.789	4.65

Table 4.1 continued from previous page

Algorithm	Noise	Trials	Accuracy (%)	Processing Time (s)
		Trial 2	0.817	4.53
		Trial 3	0.837	4.92
		Average	0.814	4.70
		Trial 1	0.728	4.91
	Misc.	Trial 2	0.715	4.89
		Trial 3	0.734	4.66
		Average	0.745	4.81

TABLE 4.1: Experimental results

4.2 Graphical Presentation of Data

Since the data above will be changed in the near future, I analyzed it much shallower and thus without a graphical representation.

## Chapter 5

# Analysis & Conclusions

### 5.1 Experimental Analysis and Trends

In the visualizations of the experiment results above, a lot can be said about to what extent MFCCs outperform DWTs, LPCs, and PLPs. As observed in Figure 4.1, Mel spectrograms vastly outperformed all of the feature extraction algorithms in terms of accuracy when there is no background noise present, but when taking time into account, MFCCs are slightly faster to process. Nonetheless, usually during practical implementations of speech recognition, unlike these trials, there will most likely be less than eight minutes of total audio data to transcribe (500 trials of 1-second audio clips), making time complexity, in this case, less important than in other programming scenarios. It also should be mentioned that the pre-processing time during model training is irrelevant to the analysis of these feature extraction algorithms, as solely the final product of these trials is assessed.

There is also a clear correlation between the performances of the spectrograms and the prediction models in various aspects. The spectrograms, including MFCCs, were much more negatively affected by the addition of background noise, which confirms the speculations of such found during my search, but still performed better than LPCs and PLPs. This could be due to the fact that prediction models such as PLP rely much more on extracting relational data instead of literal data.

The experimental process could also be improved, as it led to a few sources of error: the first of which lies in the consistency of the speed of the computer the tests were ran on. This error can be disregarded if an uncertainty of  $\pm 0.1$  is added to the running times of each trial. Another possible source of error is overfitting of the training data set from especially fast learning, which may have caused a lower than optimal performance on test datasets than could have been corrected if the model stopped training earlier. I did not account for this error since I can derive my conclusions from purely comparative data between the feature extraction algorithms.

I believe that it would be worthwhile to examine the extent to which MFCCs can perform well in noise-heavy environments with the use of noise cancellation software; raw experimental results do not draw out clear conclusions on whether MFCCs should be used if a system expects a high magnitude of background noise.

## 5.2 Further Research Opportunities and Aspirations

There are several further research opportunities and aspirations I'd like to explore in the future regarding and related to my research question. The speech feature extraction algorithms I examined in this paper are a portion of the most conventionally used computer scientists use today for speech recognition. Even so, there are some modern and advanced yet ambiguous speech feature extraction algorithms out there that may perform even better than the ones investigated in this report but don't have the recognition to gain traction in the field of ASR. Investigating the extent to which these are viable as a replacement for MFCCs could yield interesting results.

Furthermore, in the future, I wish to extensively develop this RNN model further until it's useable in practical cases. I could research the coupling of bi-directional LSTMs with these feature extraction methods (which are very popular in NLP), work with deep LSTM networks, or even experiment with Gated Recurrent Units or transforms to see how each compares.

## 5.3 Conclusions

This paper presented several conclusions regarding the research question at hand through an analysis of the performances of different speech feature extraction algorithms. The results show that although MFCCs are more optimized and contain denser key vocal features of a waveform, Mel spectrograms overall possess the highest degree of accuracy while maintaining a reasonable processing time with large inputs.

## Chapter 6

# Bibliography

A. Biswal, "Recurrent neural network (RNN) tutorial: Types and examples [updated]: Simplilearn," Simplilearn.com, 11-Aug-2022. [Online].

Available: <https://www.simplilearn.com/tutorials/deep-learning-tutorial/rnn>. [Accessed: 02-Oct-2022].

A. Graves, A.-rahman Mohamed, and G. Hinton, "Speech recognition with deep recurrent neural networks," arXiv.org, 22-Mar-2013. [Online].

Available: <https://arxiv.org/abs/1303.5778>. [Accessed: 02-Oct-2022].

A. Graves, S. Fernandez, F. Gomez, and J. Schmidhuber, "Connectionist Temporal Classification: Labelling unsegmented sequence ...," Department of Computer Science, 2006. [Online].

Available: [https://www.cs.toronto.edu/graves/icml\\_2006.pdf](https://www.cs.toronto.edu/graves/icml_2006.pdf). [Accessed: 02-Oct-2022].

A. Hannun, "Sequence Modeling with CTC," Distill, 03-Jan-2020. [Online].

Available: <https://distill.pub/2017/ctc/>. [Accessed: 02-Oct-2022].

A. Lindgren and G. Lind, "Language Classification Using Neural Networks," Uppsala Universitet, Jun-2019. [Online].

Available: <https://www.diva-portal.org/smash/get/diva2:1324945/FULLTEXT01.pdf>. [Accessed: 02-Oct-2022].

A. Taspinar, "A guide for using The wavelet transform in machine learning," ML Fundamentals, 21-Dec-2018. [Online].

Available: <https://ataspinar.com/2018/12/21/a-guide-for-using-the-wavelet-transform-in-machine-learning/>. [Accessed: 02-Oct-2022].

B. Ramakrishnan, "Lecture 16: Connectionist Temporal Classification, sequence prediction," YouTube, 02-May-2021. [Online].

Available: <https://www.youtube.com/watch?v=RowViowx1Bg>. [Accessed: 02-Oct-2022].

B. Ramakrishnan, "Lecture 17: Connectionist Temporal Classification (CTC), sequence to sequence prediction," YouTube, 02-Apr-2021. [Online].

Available: <https://www.youtube.com/watch?v=5Rj0J9AuGw0>. [Accessed: 02-Oct-2022].

C. Diatkine, "Linear Predictive coding," Introduction - Linear Predictive Coding, 2011. [Online].

Available: <https://support.ircam.fr/docs/AudioSculpt/3.0/co/LPC.html>. [Accessed: 02-Oct-2022].

H. Bidgoli, *Encyclopedia of Information Systems*. Amsterdam: Acad. Press, 2003.

H. Scheidl, "An intuitive Explanation of Connectionist Temporal Classification," Medium, 07-Jul-2021. [Online].

Available: <https://towardsdatascience.com/intuitively-understanding-connectionist-temporal-classification-3797e43a86c>. [Accessed: 02-Oct-2022].

IBM Cloud Education, "What are recurrent neural networks?," IBM, 14-Sep-2020. [Online].

Available: <https://www.ibm.com/cloud/learn/recurrent-neural-networks>. [Accessed: 02-Oct-2022].

"Introduction to the discrete wavelet transform (DWT)," Machine Intelligence Lab, 15-Feb-2004. [Online].

Available: [https://mil.ufl.edu/nechyba/www/eel6562/course\\_materials/t5.wavelets/intro\\_dwt.pdf](https://mil.ufl.edu/nechyba/www/eel6562/course_materials/t5.wavelets/intro_dwt.pdf). [Accessed: 02-Oct-2022].

J. O. Smith, "Spectral Audio Signal Processing," CCRMA, 2011. [Online].

Available: <https://ccrma.stanford.edu/jos/sasp/>. [Accessed: 02-Oct-2022].

J. S. Garofolo et al., "Timit acoustic-phonetic continuous speech corpus," Philadelphia: Linguistic Data Consortium, 1993. [Online].

Available: <https://catalog.ldc.upenn.edu/LDC93S1>. [Accessed: 02-Oct-2022].

J. Zhang, "English speech recognition system model based on computer-aided function and neural network algorithm," Computational Intelligence and Neuroscience, 22-Apr-2022. [Online].

Available: <https://www.hindawi.com/journals/cin/2022/7846877/>. [Accessed: 02-Oct-2022].

K. Daqrouq, A. R. Al-Qawasmi, K. Y. Al Azzawi, and T. A. Hilal, "Discrete wavelet transform & linear prediction coding based method for speech recognition via Neural Network," IntechOpen, 12-Sep-2011. [Online].

Available: <https://www.intechopen.com/chapters/19512>. [Accessed: 02-Oct-2022].

K. S. Rao and M. K. E., "Speech Recognition Using Articulatory and Excitation Source Features," SpringerBriefs in Speech Technology, 2017. [Online].

Available: <https://link.springer.com/content/pdf/bbm:978-3-319-49220-9/1.pdf>. [Accessed: 02-Oct-2022].

L. L. Thomala, "Baidu: R&D spending 2021," Statista, 01-Apr-2022. [Online].

Available: <https://www.statista.com/statistics/1079978/china-baidu-research-and-development-costs/>. [Accessed: 02-Oct-2022].

L. Li, "Performance analysis of objective speech quality measures in Mel domain," Science Alert, 03-Jan-2015. [Online].

Available: <https://scialert.net/fulltext/?doi=jse.2015.350.361>. [Accessed: 02-Oct-2022].

N. S. Nehe and R. S. Holambe, "DWT and LPC based feature extraction methods for Isolated

word recognition - eurasip journal on audio, speech, and music processing,” SpringerOpen, 30-Jan-2012. [Online].

Available: <https://asmp-erasipjournals.springeropen.com/articles/10.1186/1687-4722-2012-7>. [Accessed: 02-Oct-2022].

P. Koehn, Neural machine translation. Cambridge, ON: Cambridge University Press, 2020.

P. V. Janse, S. B. Magre, P. K. Kurzekar, and R. R. Deshmukh, “A comparative study between MFCC and DWT feature extraction technique,” International Journal of Engineering Research & Technology, 30-Jan-2014. [Online].

Available: <https://www.ijert.org/a-comparative-study-between-mfcc-and-dwt-feature-extraction-technique>. [Accessed: 02-Oct-2022].

R. M. Gray, “Linear predictive coding and the Internet Protocol,” The Essence of Knowledge, 2010. [Online]. Available: <https://ee.stanford.edu/gray/lpcip.pdf>. [Accessed: 02-Oct-2022].

R. Wang, “Automatic speech recognition 101: How ASR works,” Dialpad. [Online]. Available: <https://www.dialpad.com/blog/automatic-speech-recognition/>. [Accessed: 02-Oct-2022].

S. A. Alim and N. K. A. Rashid, “Some commonly used speech feature extraction algorithms,” IntechOpen, 12-Dec-2018. [Online].

Available: <https://www.intechopen.com/chapters/63970>. [Accessed: 02-Oct-2022].

S. M, “Let’s understand the problems with recurrent neural networks,” Analytics Vidhya, 10-Jul-2021. [Online].

Available: <https://www.analyticsvidhya.com/blog/2021/07/lets-understand-the-problems-with-recurrent-neural-networks/>. [Accessed: 02-Oct-2022].

Saranga-K-Mahanta-google and arvindpdmn, “Audio Feature Extraction,” Devopedia, 23-May-2021. [Online].

Available: <https://devopedia.org/audio-feature-extraction>. [Accessed: 02-Oct-2022].

Sonix Authors, “A brief history of speech recognition,” Sonix, 2022. [Online].

Available: <https://sonix.ai/history-of-speech-recognition>. [Accessed: 02-Oct-2022].

Summa Linguae Authors, “Speech recognition software: History, present, and future,” Summa Linguae, 18-Jun-2021. [Online].

Available: <https://summalinguae.com/language-technology/speech-recognition-software-history-future/>. [Accessed: 02-Oct-2022].

V. Lendave, “LSTM vs gru in recurrent neural network: A comparative study,” Analytics India Magazine, 28-Aug-2021. [Online].

Available: <https://analyticsindiamag.com/lstm-vs-gru-in-recurrent-neural-network-a-comparative-study/>. [Accessed: 02-Oct-2022].

Y. Wang, X. Deng, S. Pu, and Z. Huang, “Residual convolutional CTC networks for automatic speech recognition,” arXiv.org, 24-Feb-2017. [Online].

Available: <https://arxiv.org/abs/1702.07793>. [Accessed: 02-Oct-2022].



## **Appendix A**

# **Appendix A**

### **A.1 Source Code**

The source code will be uploaded to a public GitHub repository associated to the profile "FinityFly"